

Dynamic Multi-Raft

Dongxu Huang
PingCAP



CNUTCon 2017

全球运维技术大会

上海·光大会展中心大酒店 | 2017.9.10-11

智能时代的新运维

大数据运维
安全
SRE
Kubernetes
Serverless
游戏运维
智能化运维
基础架构
监控
DevOps
AIOps
互联网金融





斯达克学院

实践驱动的IT教育



斯达克学院(StuQ)，极客邦旗下实践驱动的IT教育平台。通过线下和线上多种形式的综合学习解决方案，帮助IT从业者和研发团队提升技能水平。



10大职业技术领域课程

<http://www.stuq.org>

SPEAKER INTRODUCE

黄东旭 CTO & Cofounder, PingCAP

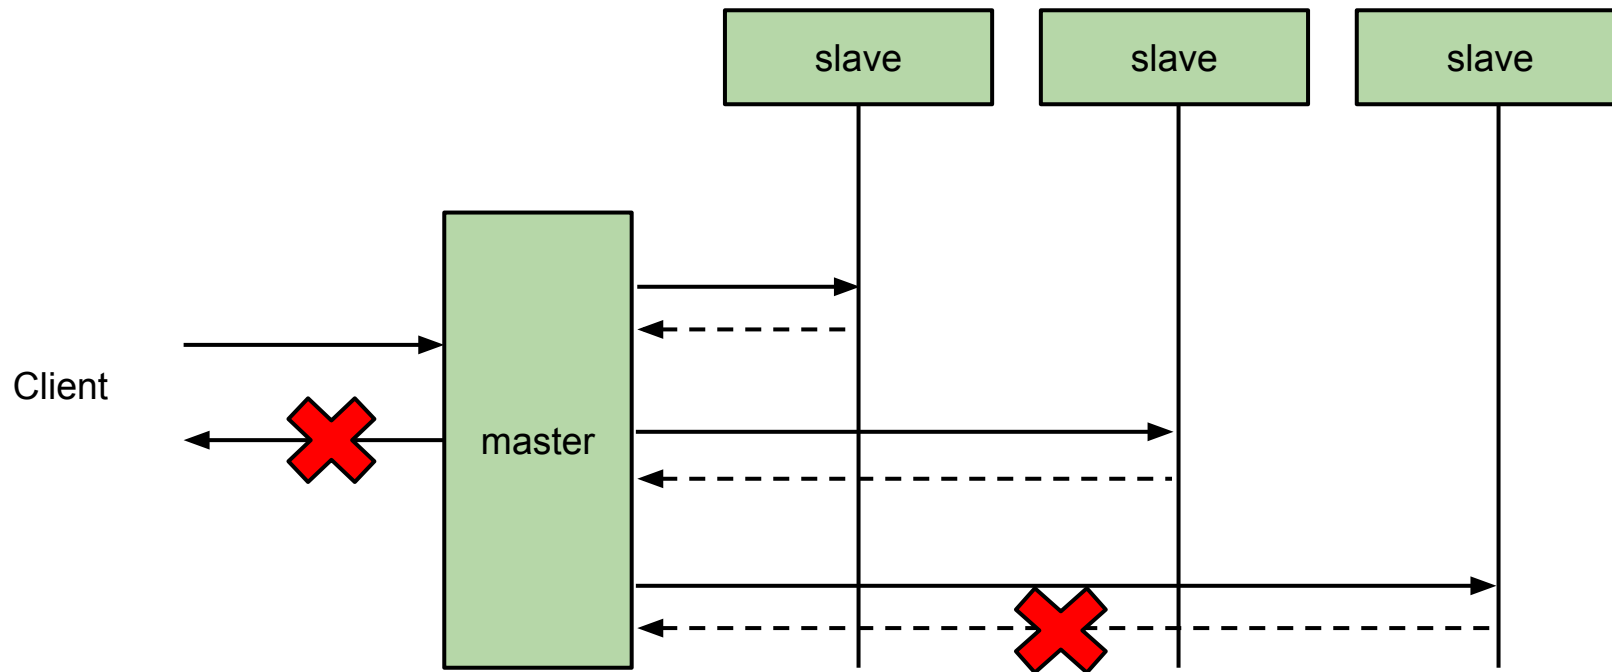
- MSRA / Netease / WandouLabs / PingCAP
- Hacker / Infrastructure software engineer
- Distributed system / Database / PL / ...
- Codis / TiDB / TiKV
- Golang / Rust / Python

Consensus

is the only problem in distributed system...

Modern HA

- Master-slave is not an option, why?



MySQL MHA + Semi-Sync?

Bug #80395		semi-sync: incorrect crash recovery handling	
Submitted:	16 Feb 2016 12:54	Modified:	16 Feb 2016 18:17
Reporter:	Matt Lord	Email Updates:	Subscribe
Status:	Need Doc Info	Impact on me:	None Affects Me
Category:	MySQL Server: Replication	Severity:	S2 (Serious)
Version:	5.7.11	OS:	Any
Assigned to:	David Moss		
<div> View Add Comment Files Developer Edit Submission View Progress Log Contributions </div>			

[16 Feb 2016 12:54] Matt Lord

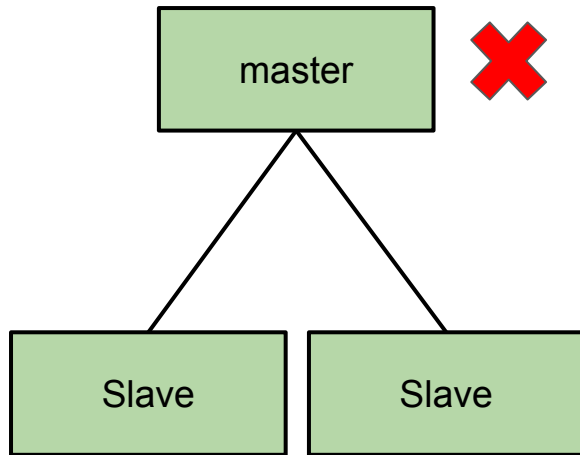
Description:

When mysqld is killed while an open semi-sync replication transaction is waiting for the master timeout, that prepared *but uncommitted* transaction is NOT property rolled back when the master performs its subsequent automated crash recovery.

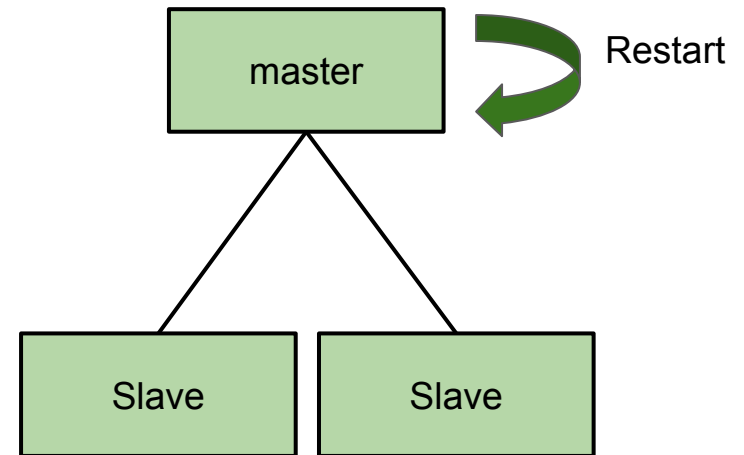
This was verified on OL 7.2 x86_64, using MySQL 5.7.11-community.

How to repeat:

(1) Insert A, but master crash when doing semi-sync

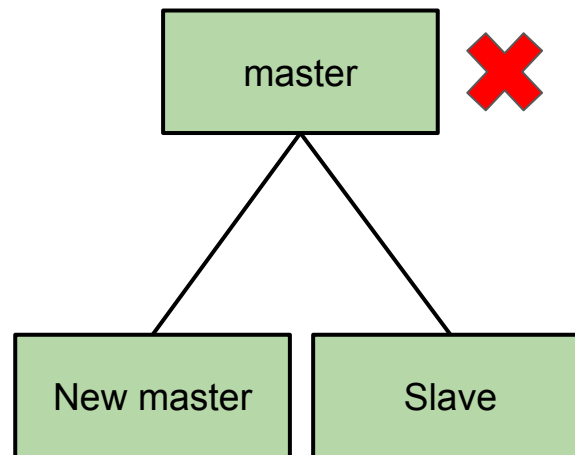


(2) Read A, OK; But degrade to async-replication



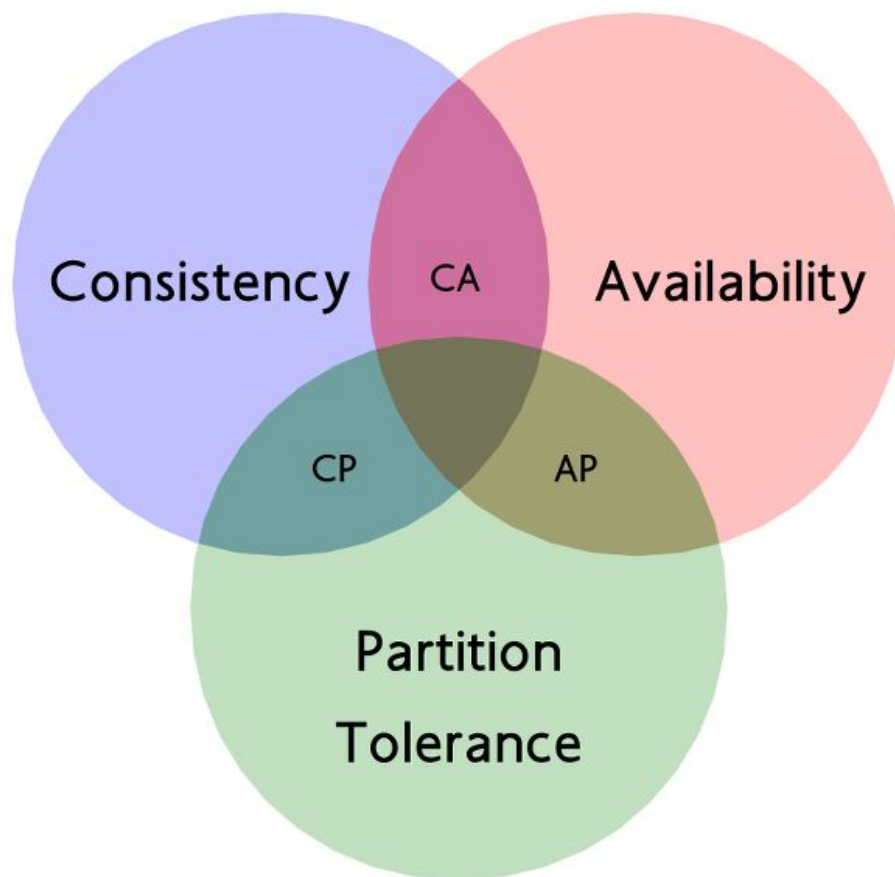
(3) Master crash again...then new master is elected

Read A, fail, oops...

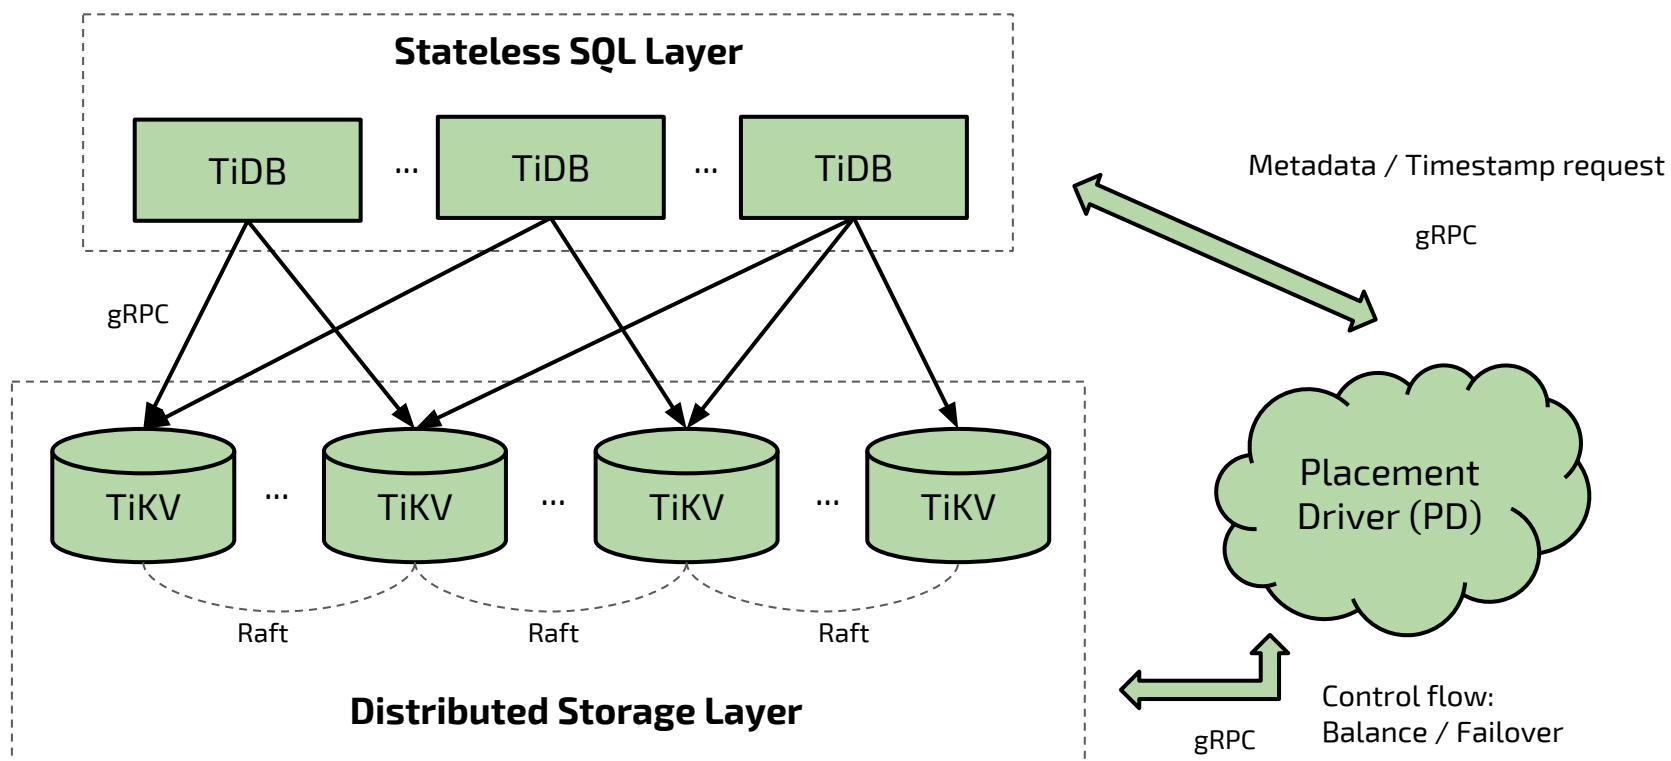


TiDB Project (Requirement)

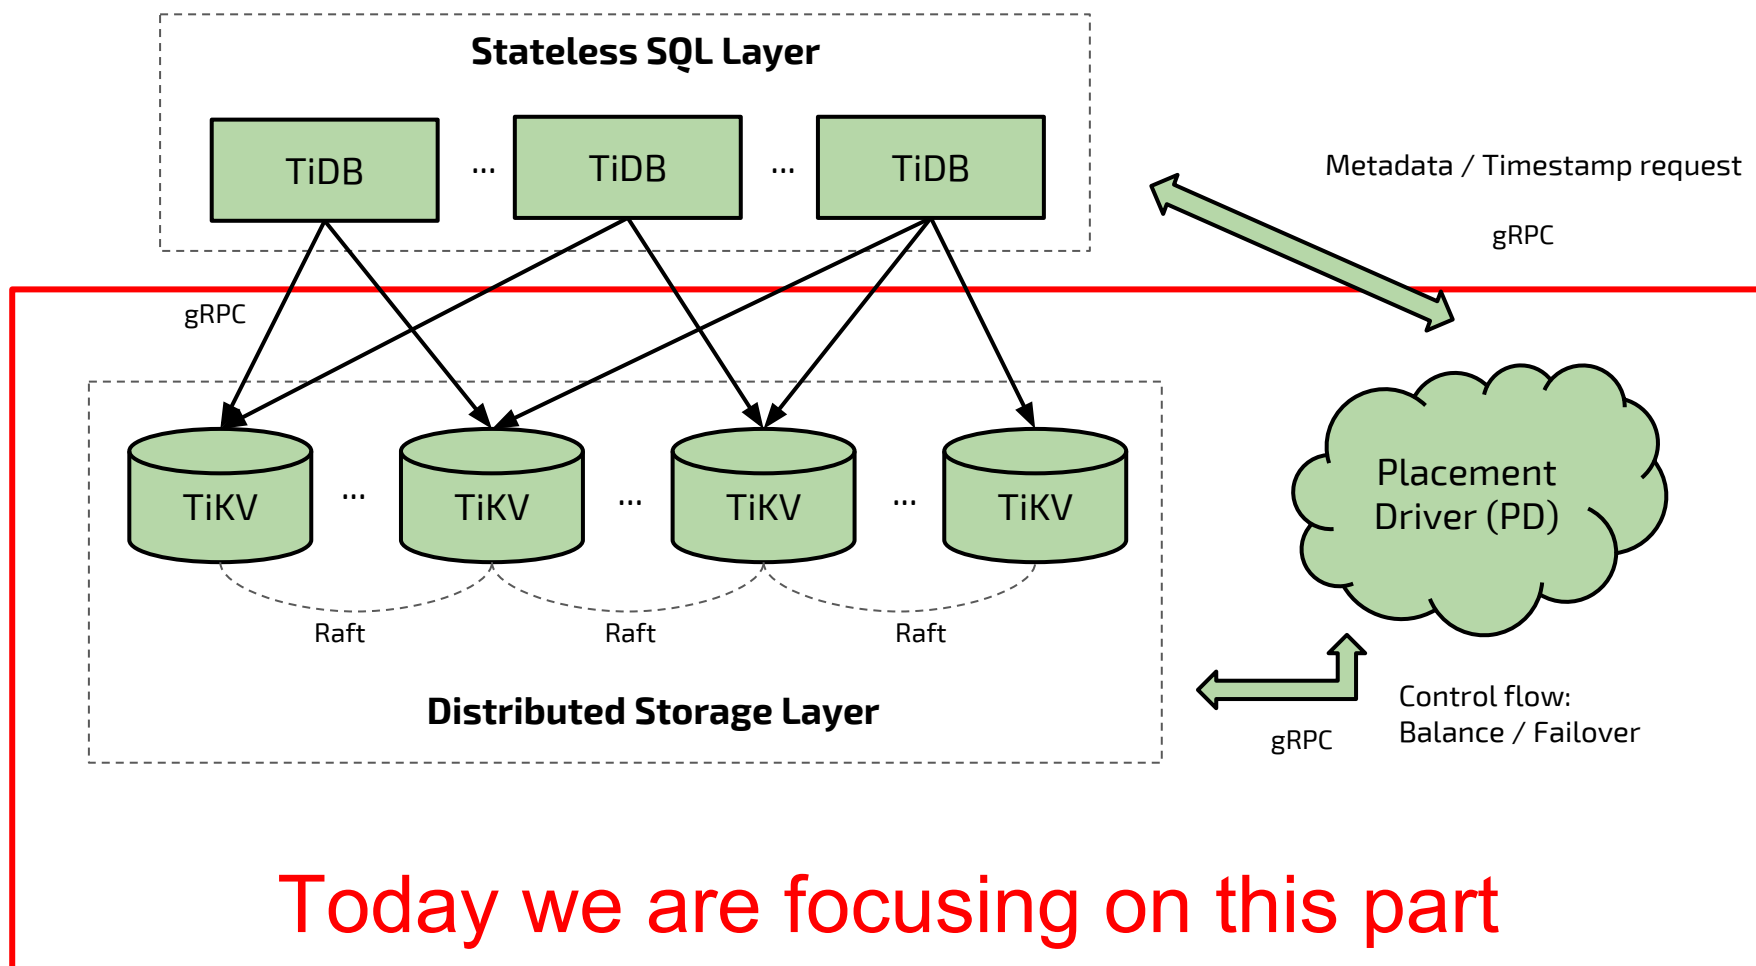
- Strong consistency
- Scalability
- High availability



TiDB Project Overview



TiDB Project Overview



Replicated State Machines

- All servers execute same commands in same order
- System makes progress as long as any majority of servers up
- Agreement on shared state (single system image)
- Recovers from server failures autonomously
 - Minority of servers fail: no problem
 - Majority fail: lose availability, retain strong consistency
- IMHO, there are only two RSM implementations:
 - Multi-Paxos / Raft

The problem in Paxos (Multi Paxos)

“The dirty little secret of the NSDI community is that at most five people really, truly understand every part of Paxos;-).”

—NSDI reviewer

Raft saves the day

- **Leader election**

- Select one of the servers to act as cluster leader
- Detect crashes, choose new leader

- **Log replication**

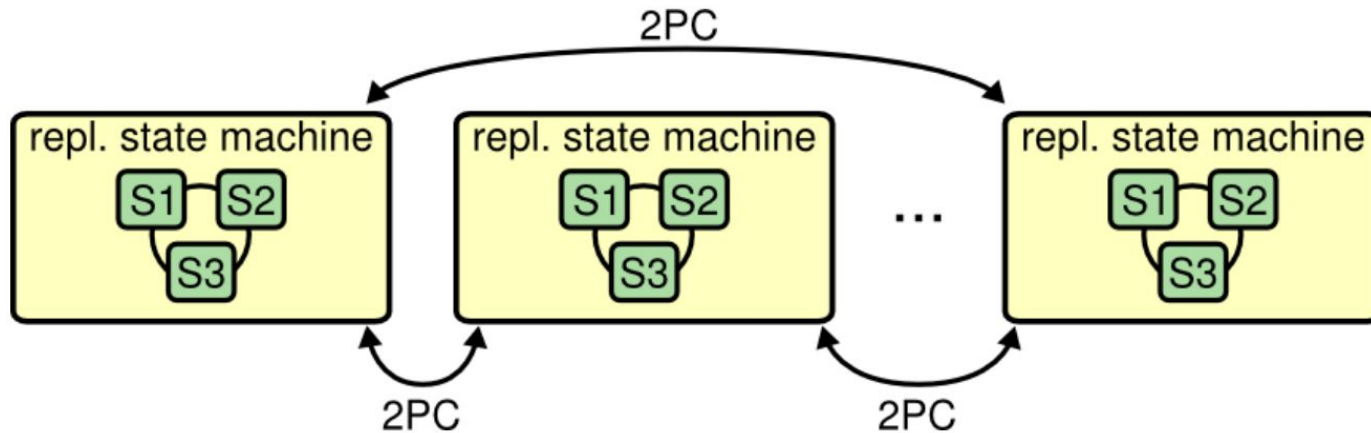
- Leader takes commands from clients, appends to its log
- Leader replicates its log to other servers (overwriting inconsistencies)

- **Safety**

- Only a server with an up-to-date log can become leader

Use Raft in database

- Single RSM is **NOT** gonna work.
- You need 2PC to retain strong consistency across different RSMs.



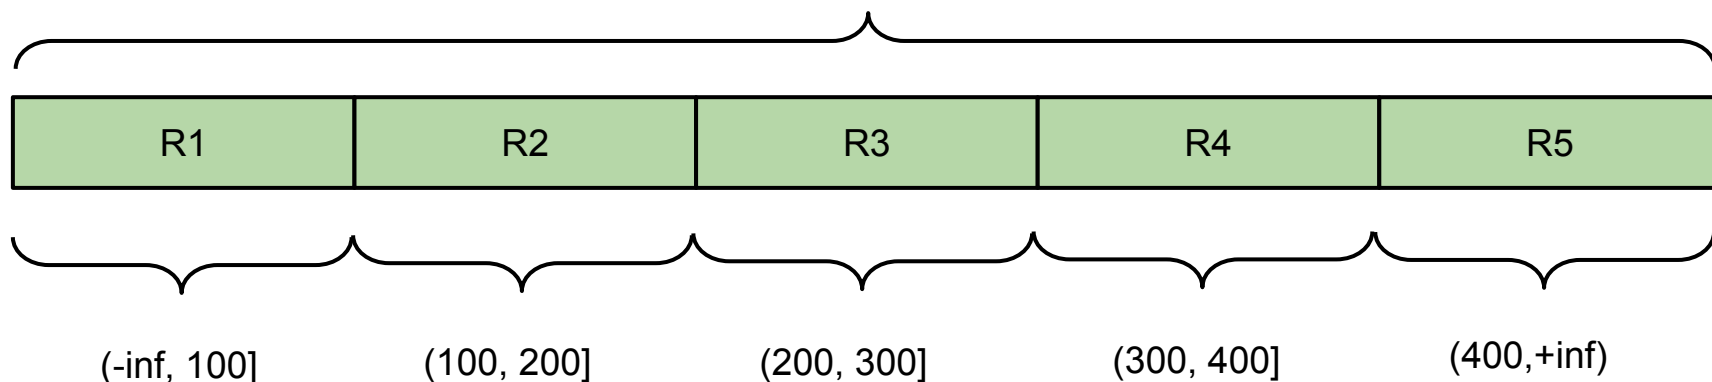
Raft in database

- How to shard?
- How to split / merge dynamically?
- How to balance the workload?
- How to improve the throughput?

Sharding Raft in TiKV

- Split key space into **Regions (normally in byte-order) logically**
- Each region is a raft group
 - Default size: 96 ~ 128 MB
 - Why?

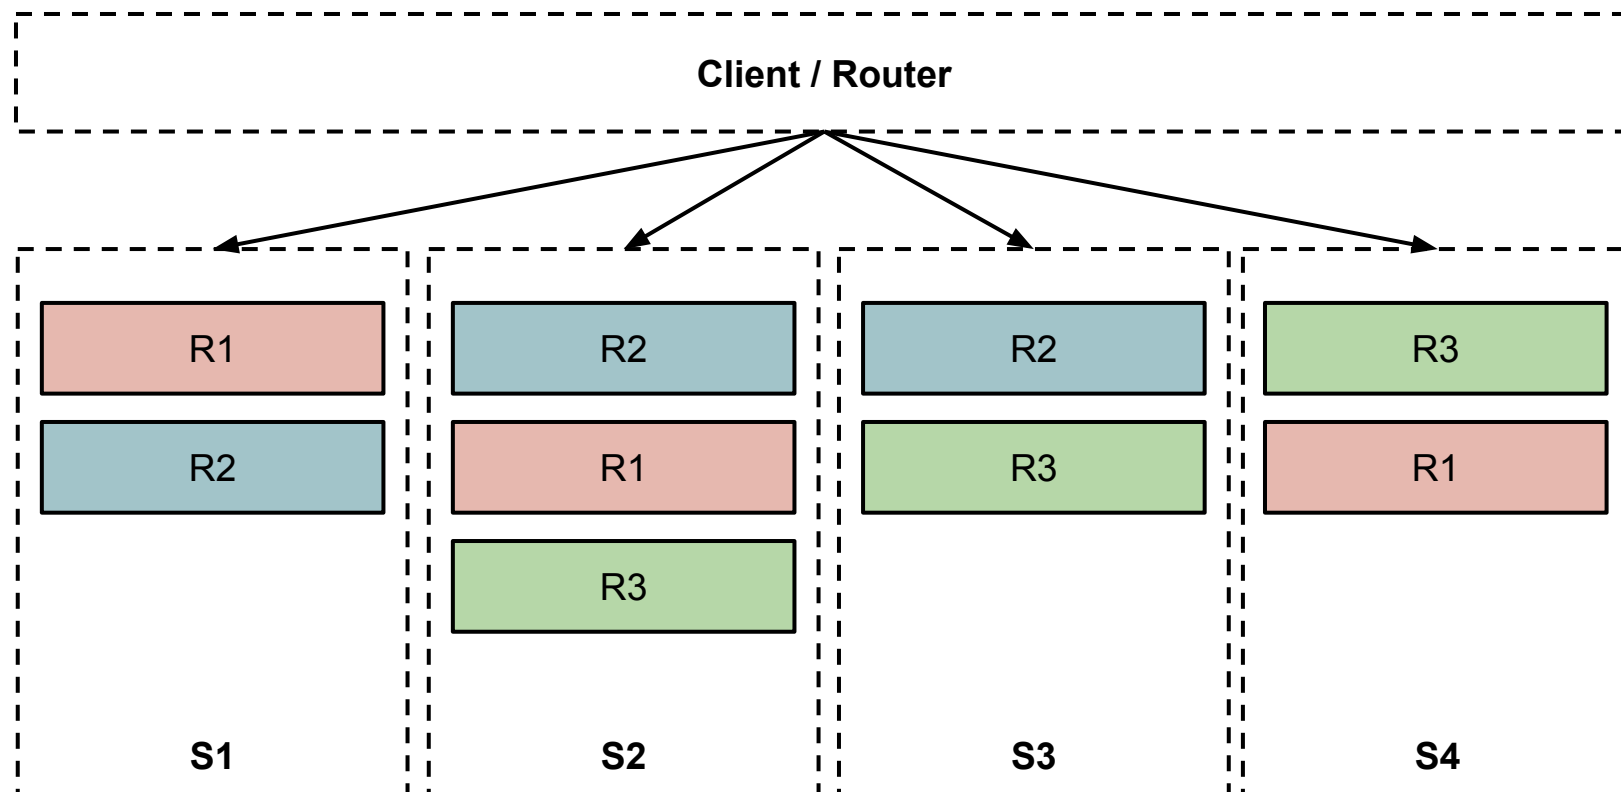
Key space (-inf, +inf)



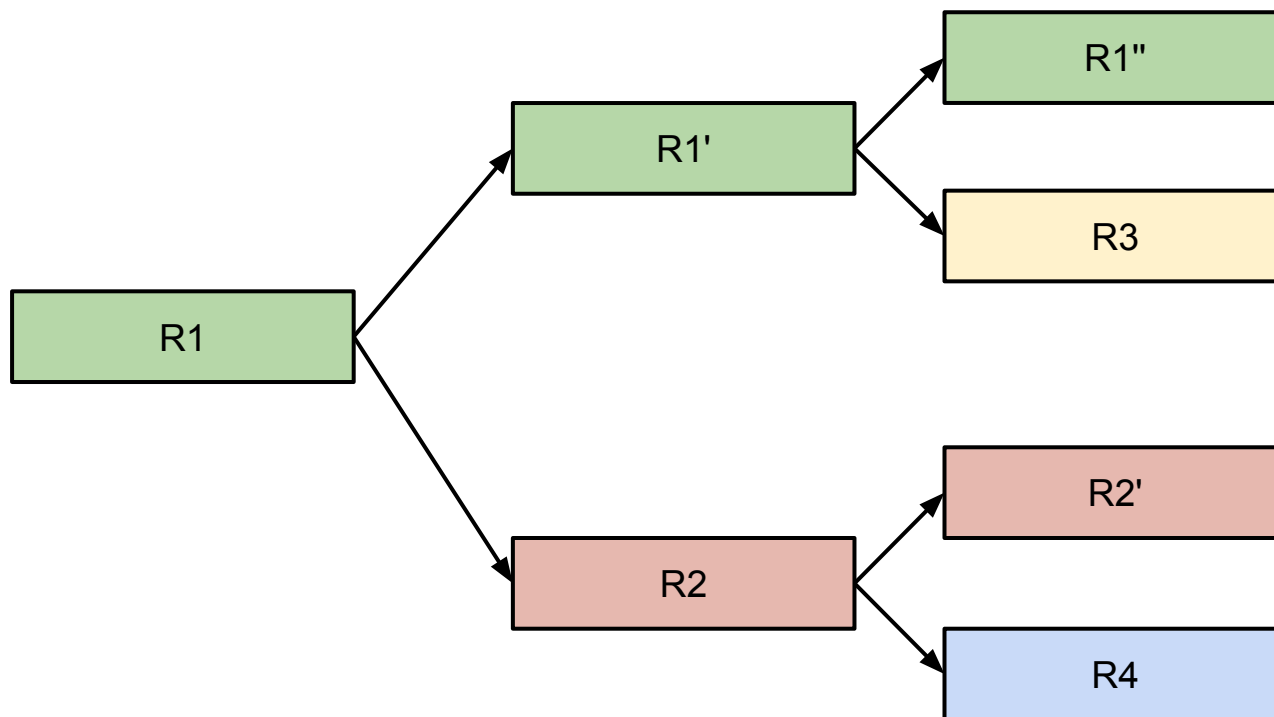
Meta data storage

- We store region meta in an in-memory B-Tree (in PD)
 - Sorted by the start key of region
 - We can find the right region which contains specific key in $O(\log N)$
- PD is not '***the source of truth***', data server is. Why?
 - Split is always happening
 - The metadata stored in PD may be out-of-date
 - Retry is important

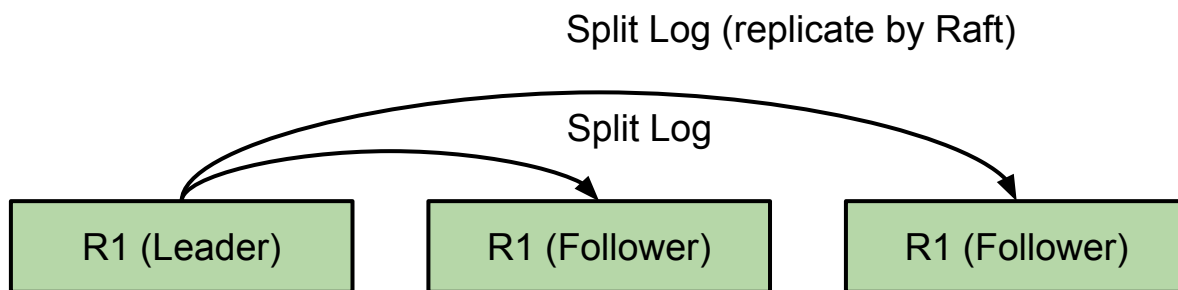
Sharding Raft in TiKV



Sharding Raft in TiKV



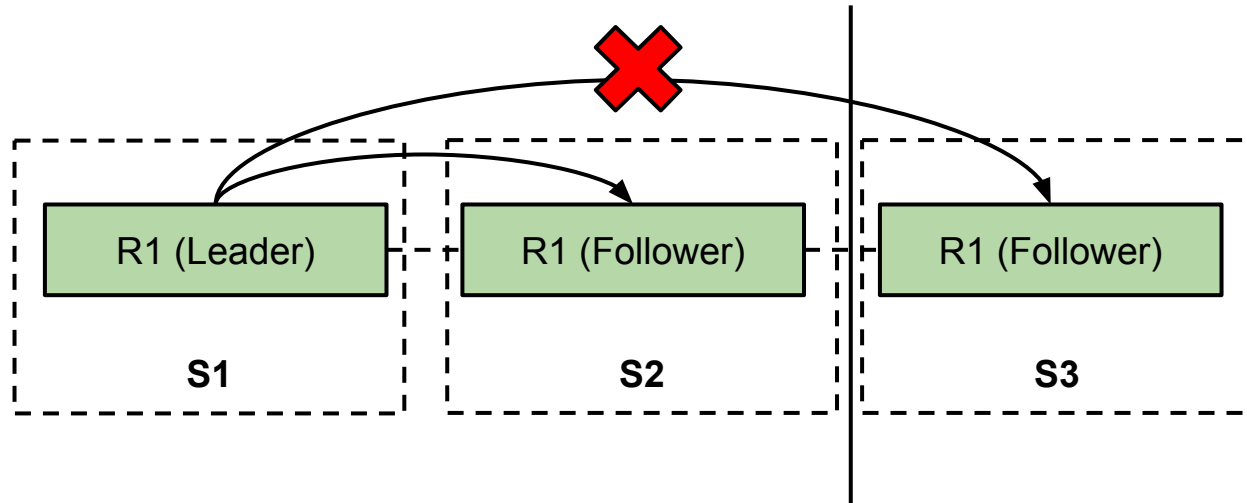
Dynamic split / merge



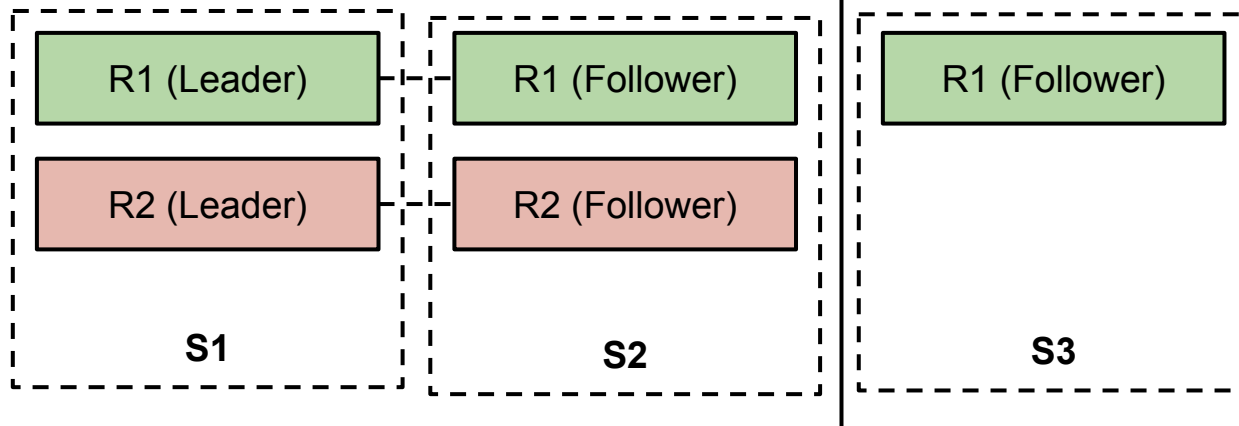
Simple...Huh?

An abnormal situation...

(1)

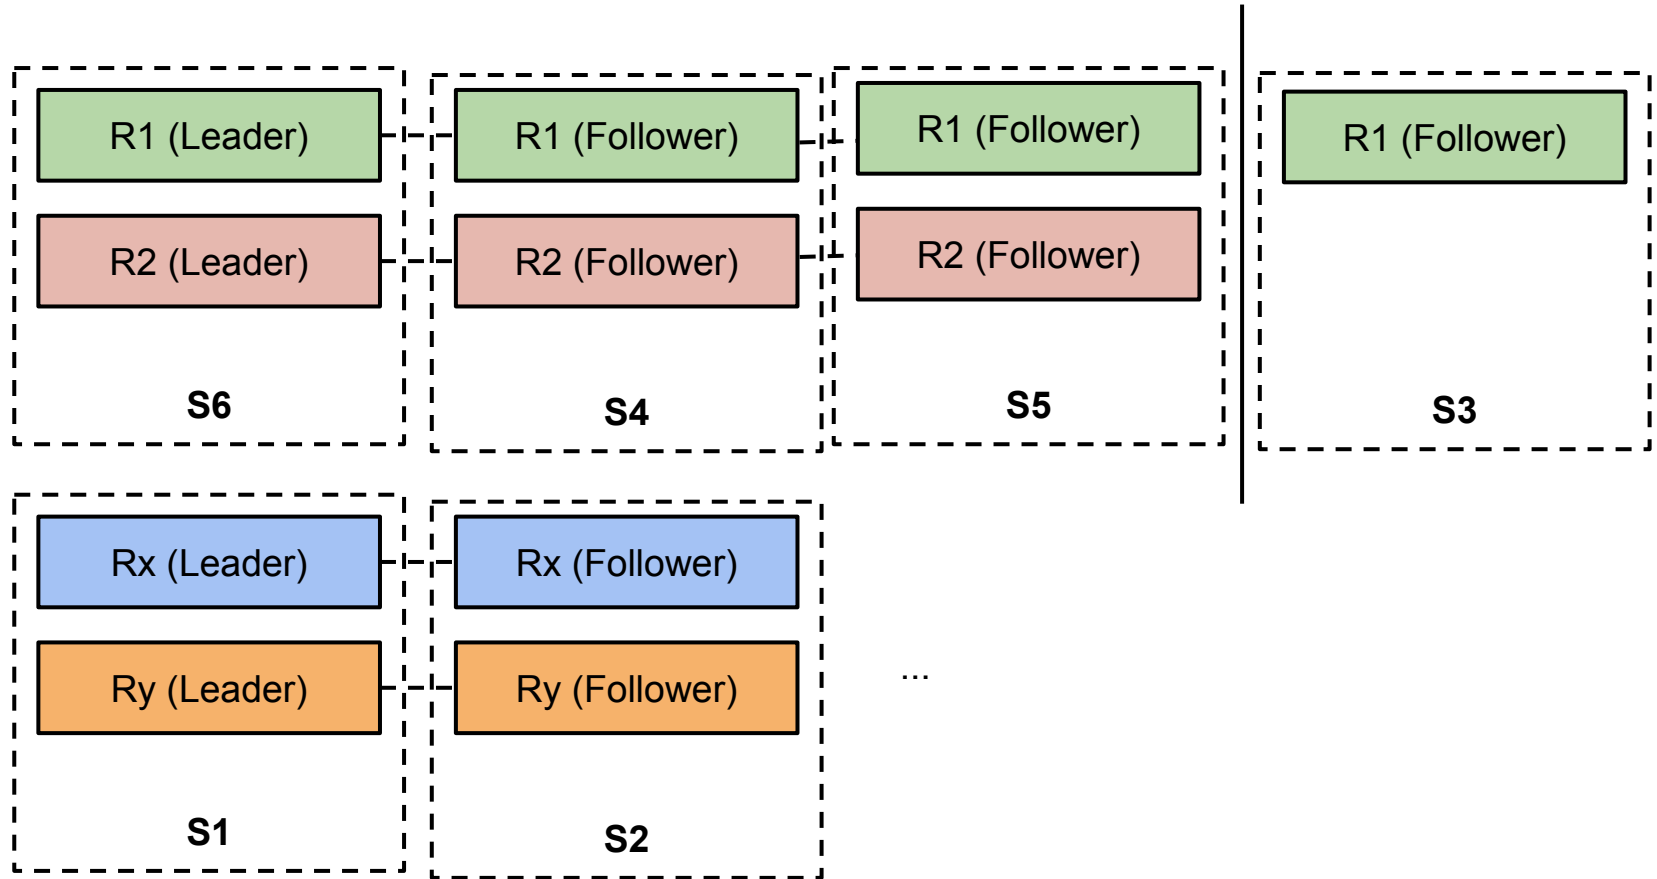


(2)



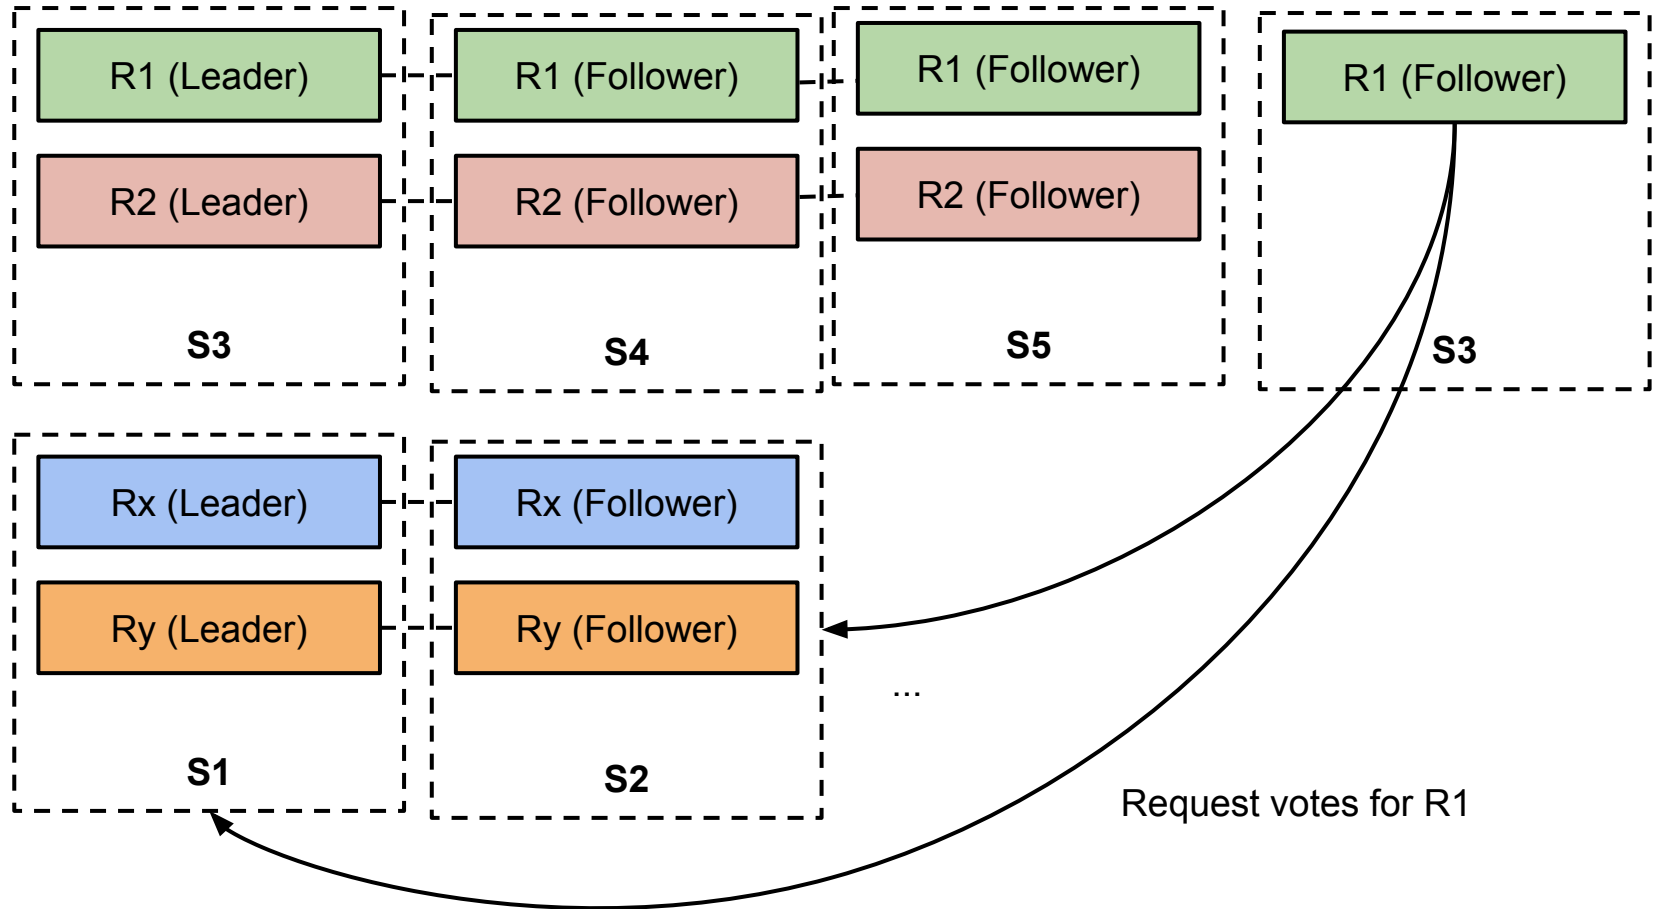
An abnormal situation...

(3) After N rounds of split or membership changes...

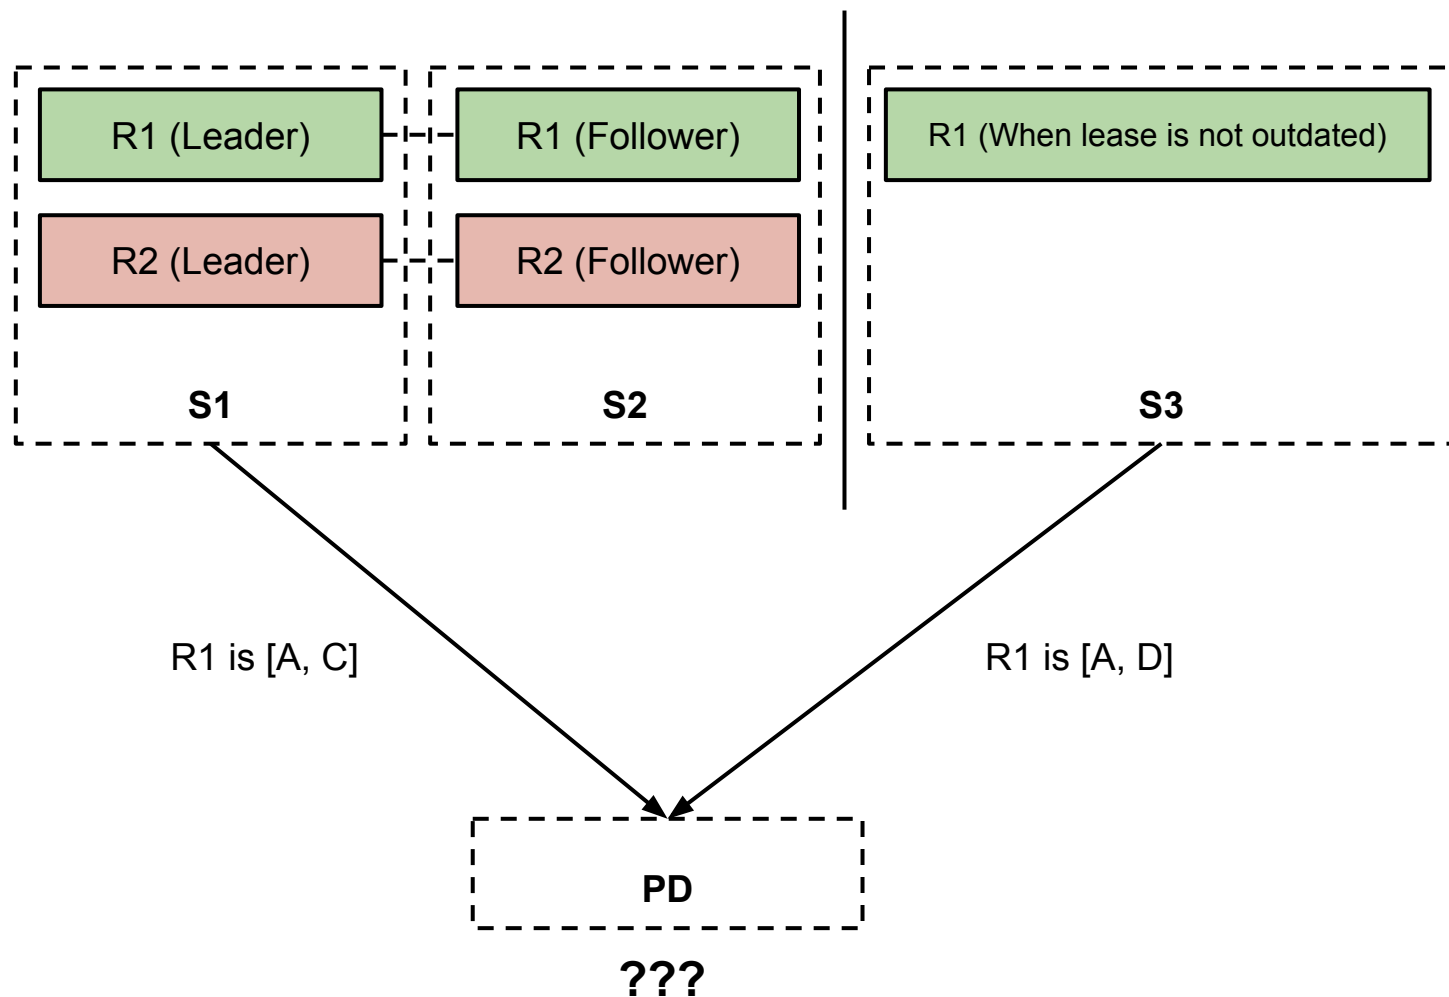


An abnormal situation...

(4)



Another abnormal situation

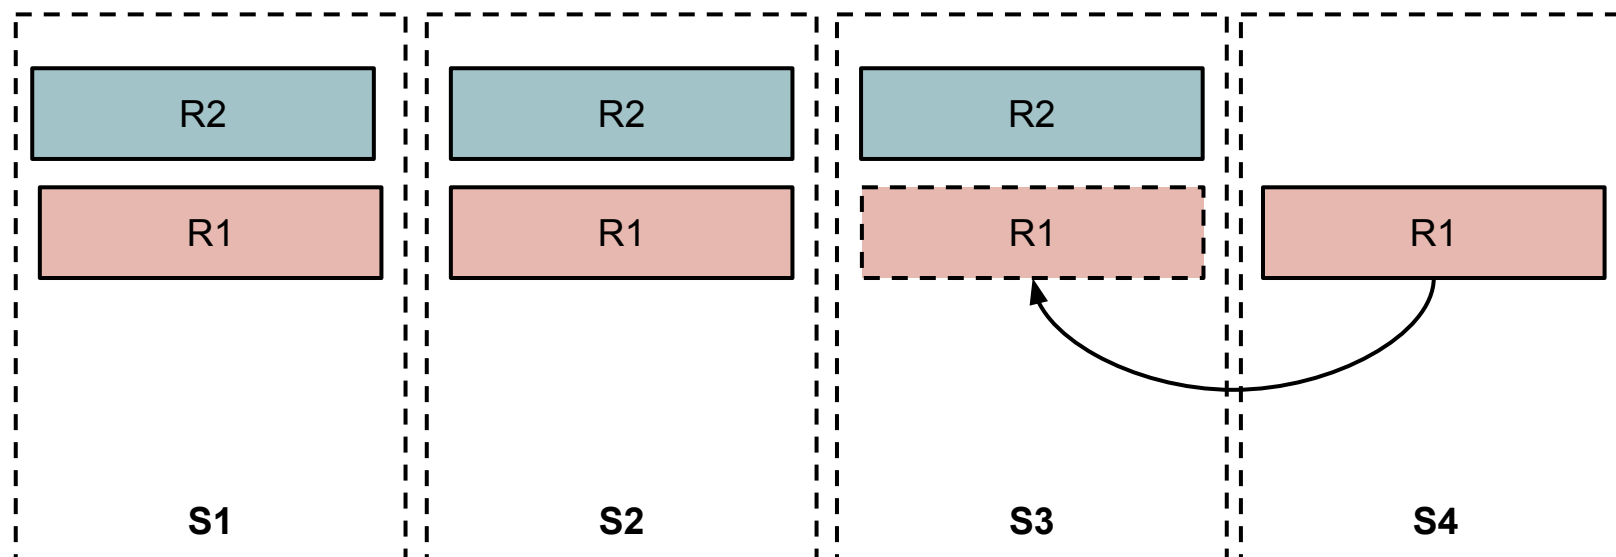


Introduce Region Epoch

- $\text{Epoch}(\text{Region } X) := \{\text{ConfVer}, \text{SplitVer}\}$
- Every configuration change in Region X will increase the ConfVer
- Every split occurs in Region X will increase the SplitVer
- Let's say $\text{Epoch}(R1) \geq \text{Epoch}(R2)$, if and only if:
 - $\text{ConfVer}(R1) \geq \text{ConfVer}(R2)$ and $\text{SplitVer}(R1) \geq \text{SplitVer}(R2)$
- Larger epoch always win

What about merge?

- Make sure all replica for these two regions are in same nodes
- And no more rebalance for these two regions

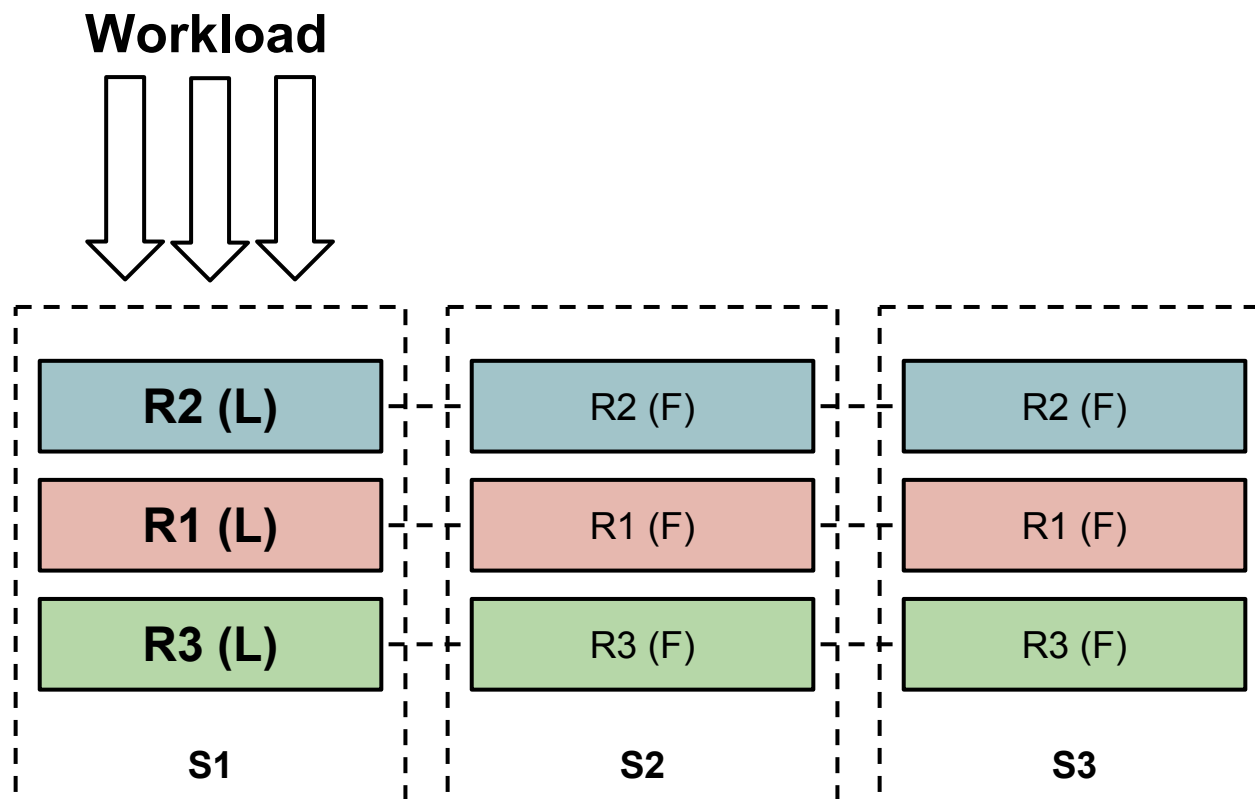


Storage

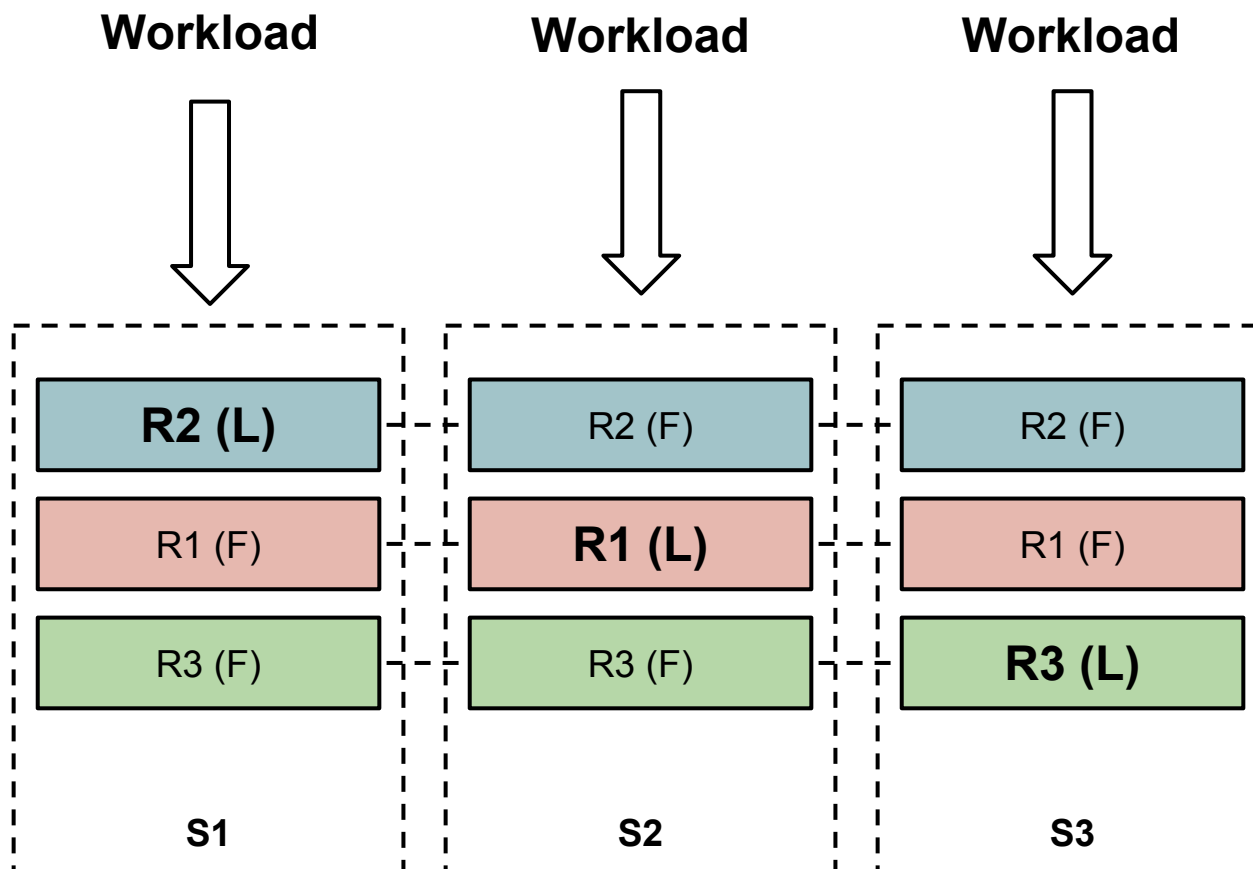
- RSM storage
 - All regions in the same physical node share one RocksDB instance
- Log storage
 - Journal-like storage
 - Share with region storage

Leadership transfer

- For fast rebalance, since Raft is a randomized algorithm, there is a certain probability that one node has too many leaders.

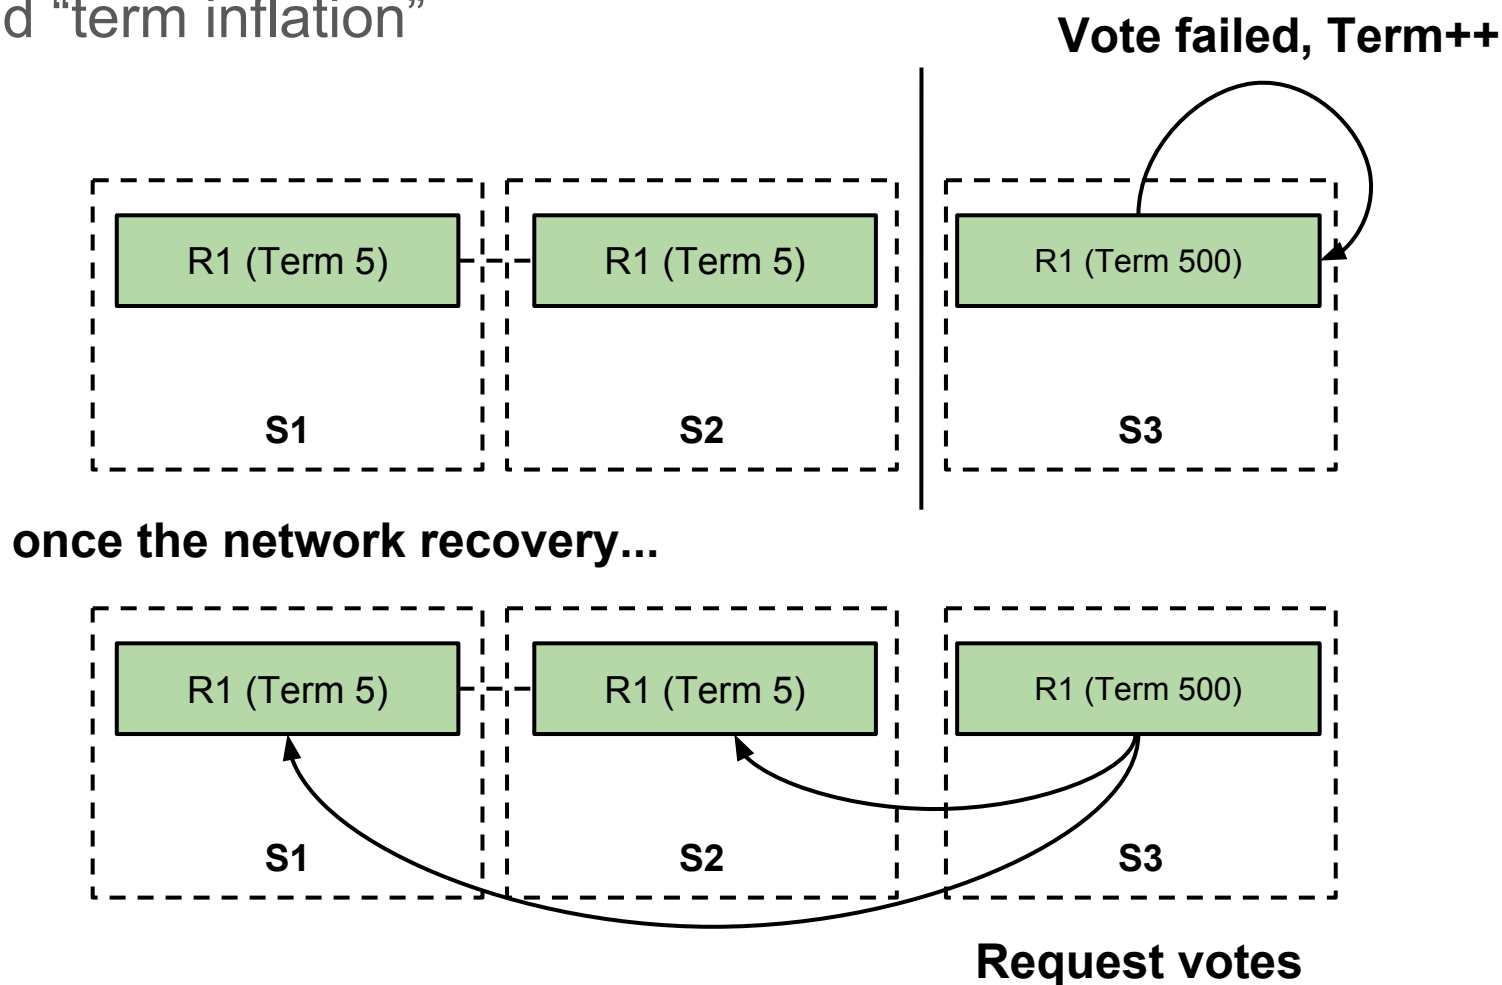


Leadership transfer

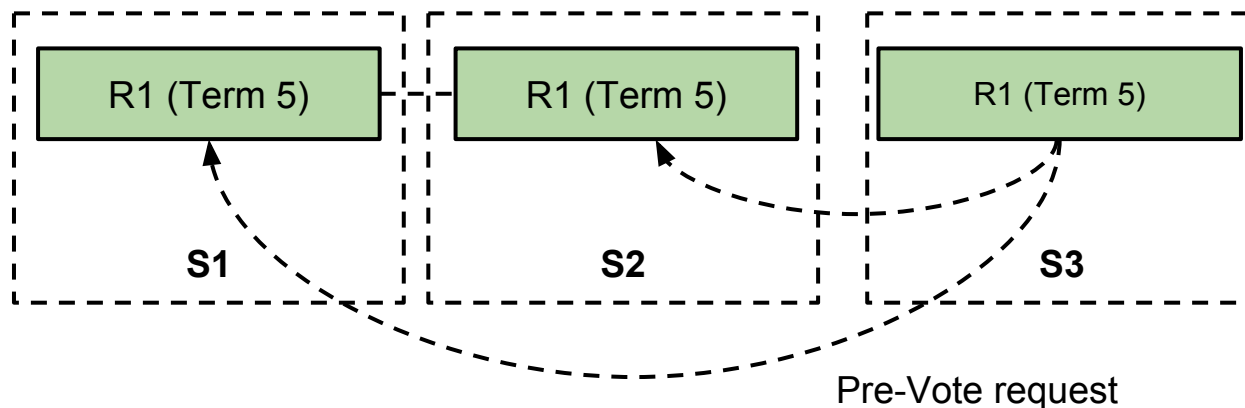


Pre-vote algorithm

- Avoid “term inflation”



Pre-vote algorithm



S3 sends Pre-vote request to S1 and S2 to make sure S3's log is up-to-date, when S3 receives responses from a majority of the cluster, S3 will increase its term and start a normal election

How to test

- Testing in distributed system is really hard
- Test-Driven Development
- Test cases from community
 - Lots of tests in MySQL drivers/connectors
 - Lots of ORMs
 - Lots of applications (Record---replay)
- Fault injection
 - Hardware: disk error, network card, cpu, clock
 - Software: file system, network and protocol
- Simulate everything : Network
- Distribute testing
 - Jepsen
 - Namazu

Benchmark

- 46 Physical nodes
- 460 TiKV instances (1 tikv instance for 1 HDD)
- TiKV Raw API Put (Raft)

Put(key, value)

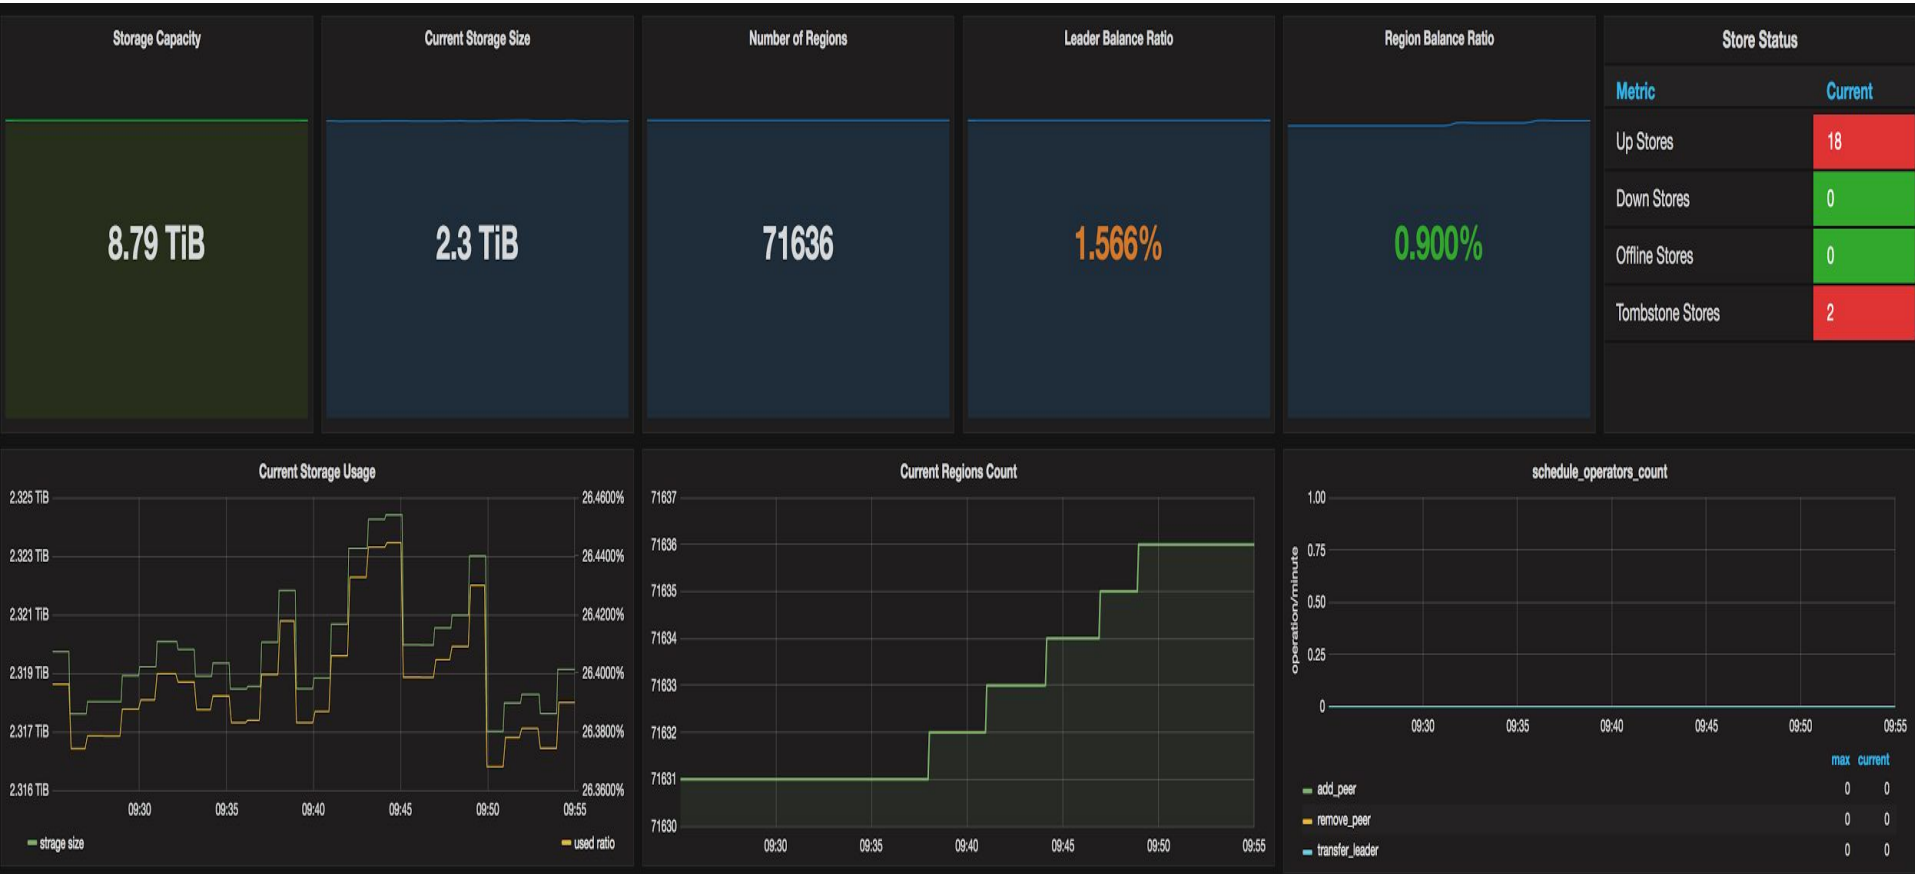
key size: 21 bytes

value size: random (1~100 bytes)

Benchmark



Benchmark



THANKS!



2017 ArchSummit 深圳
TiDB 讨论



让创新技术推动社会进步

HELP TO BUILD A BETTER SOCIETY WITH
INNOVATIVE TECHNOLOGIES

Geekbang>

极客邦科技

InfoQ

专注中高端技术人员的技术媒体



EGO EXTRA GEEKS' ORGANIZATION
NETWORKS

高端技术人员学习型社交平台



StuQ
斯达克学院

实践驱动的 IT 教育平台

